BHANU TEJA PANGULURI

(443)-497-4375 | bhanuteja.panguluri@gmail.com | [LinkedIn](#) | [Portfolio](#) | [Github](#)

EDUCATION

University of Maryland   College Park, MD
Master of Engineering in Software Engineering   *May 2025*

CERTIFICATIONS

AWS Certified Cloud Practitioner **AWS CCP - ca7a26887ca940849ede7b02a2d120d5**
CodePath Advanced Technical Interview Prep **CERT ID - 97614**
Machine Learning with Big Data **CERT ID - 3ETNE3Y62HAV**

WORK EXPERIENCE

Software Engineer   August 2023 - Present
University of Maryland - Engineering Information Technology EIT   *Maryland, USA*
- Led operations across the engineering IT portfolio and shipped targeted automations that eliminated reconciliation bottlenecks and produced consistently on-time financial reporting.
- Built an inventory tracking platform used by 10 engineering teams that unified asset lifecycle data, prevented stockouts, and improved forecasting for procurement.
- Designed normalized data models and SQL/Tableau dashboards for budget oversight that exposed spend variances and enabled 12% cost savings across departments.
- Automated audit and data-quality workflows in Python with pandas and NumPy that enforced validation rules and lifted submission accuracy to 95% for compliance.
- Implemented scheduled ETL in SQL and Python that integrated finance, purchasing, and ticketing systems, standardized KPIs, and shortened the month-end close by 2 days.
- Established structured logging and request tracing with centralized dashboards that surfaced error spikes and cut time to isolate root causes during incidents.

Data Engineer   August 2021 - July 2023
ZS Associates   *Pune, India*
- Orchestrated the Dynamic Targeting pilot for a top healthcare client and guided scale-out to multiple business units, unlocking $2M+ in follow-on revenue and executive sponsorship.
- Engineered production ETL with Python, SQL, Hive, and PySpark on AWS that processed multi-terabyte workloads and cut batch runtime by 40% against the initial baseline.
- Built reusable preprocessing pipelines in Python and PySpark that handled feature engineering, data cleansing, and HCP identity resolution. Reduced data preparation time by 35% and improved model stability across retrains.
- Analyzed 250+ user feedback responses and ran A/B tests to improve feature discoverability, raising product adoption from 40% to 70% within 2 quarters.
- Drove propensity modeling and exploratory analysis that prioritized HCP outreach. Improved rep productivity by 16% and increased conversion on targeted segments by 12% versus baseline.
- Delivered weekly KPI suites in Tableau and Power BI backed by Hive that gave leaders clear visibility into adoption and impact, and helped win the full-scale project

Software Development Engineer   January 2021 - July 2021
Toshiba   *Bangalore, India*
- Developed C++ and CUDA backend modules for MRI and CT image reconstruction that increased reconstruction throughput by 22% and met real-time latency targets.
- Optimized detection pipelines and trained models in Python that improved sensitivity by 8% and reduced false positives by 10% on internal benchmark datasets.
- Tuned GPU kernels with shared memory and asynchronous streams that delivered 1.9× speedup on compute heavy routines and lowered device memory usage by 15%.
- Expanded automated unit and integration test suites that raised code coverage from 62% to 86% and cut regression defects in staging by 30%.

SKILLS

Languages: Java, Python, C/C++, Node.js, JavaScript, Scala, PHP, Bash
Platforms: PostgreSQL, MySQL, SQL Server, MongoDB, DynamoDB, Redis, Elasticsearch, Snowflake, BigQuery, Redshift, Azure Synapse, Databricks, Amazon S3, Azure Data Lake
Pipelines: Apache Spark, PySpark, Kafka, Apache Airflow, dbt, AWS Glue, Azure Data Factory, Google Dataflow, Hadoop, Hive
Analytics: Power BI, Tableau, Exploratory Analysis, AB Testing, Time Series, Regression, Forecasting
Architecture & Ops: REST, GraphQL, FastAPI, OAuth2, JWT, Git, CI/CD, Docker, Kubernetes, Terraform, Linux, OpenTelemetry, Prometheus, Grafana, Datadog, AWS, Azure, GCP